

CookBook Report

Project Team Albertosaurus:

Bhairaw Aryan

Jeremy Ng

Rushil Punchhi

Ting Xu

Submitted to:

Professor Mieszko Lis

CPEN 291

Abstract

CookBook is a web and android app that takes in images of fruits and vegetables, uses a machine learning model to recognize them in the images, and then outputs various recipes that use them as ingredients. Users can either upload images from their computer, take an image straight from their camera, or type in the name of an ingredient. Each recipe's anatomy can further be fetched to display the ingredients, instructions, prep time, serving size, wine-pairing, cuisine type and a link to the original source of the recipe. The latter is achieved by using the Spoonacular client — an API that connects over 360,000 recipes as well as an open source recipe database for complex food ontology.

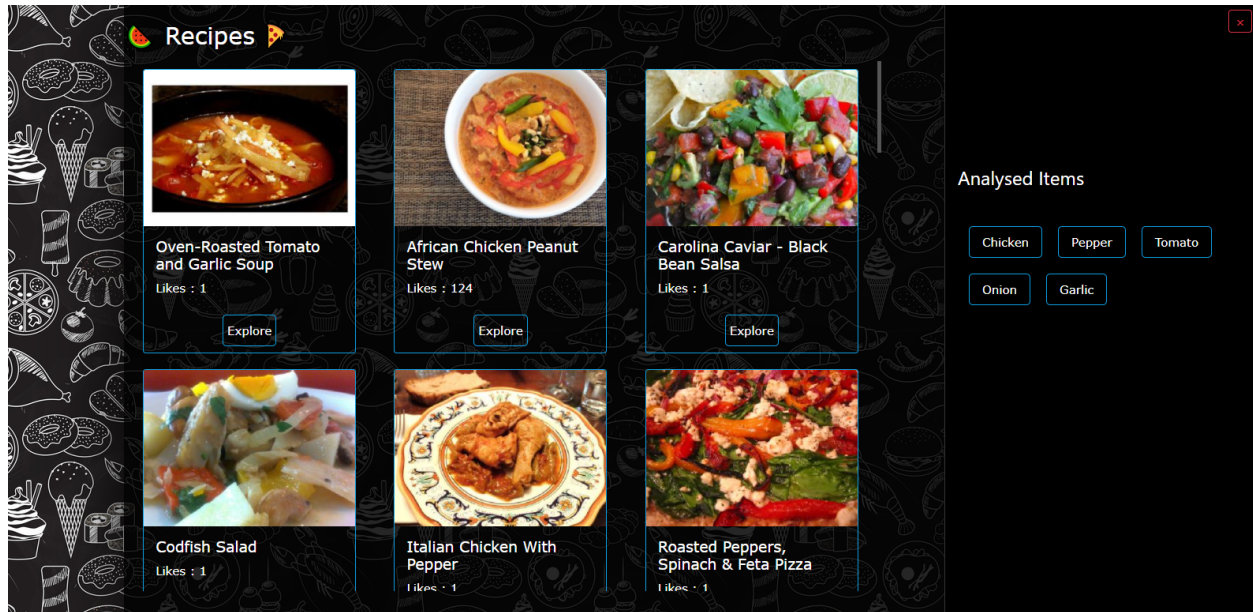
Walkthrough

When first connecting to CookBook via the web app, users will see a home screen with a 'Get Started!' button. The button pops an interface, with the option to either upload images of the ingredients through camera(if present)/regular file upload or manually type in the name of the ingredients. In the image below, we have typed in 'Pepper', 'Chicken' and uploaded images of garlic, onions, and tomatoes.



Main page of CookBook.

After pressing upload, CookBook will use its machine learning model to recognize the fruits or vegetables in the image, and then use the Spoonacular API to fetch suitable recipes. In the image below, the identified ingredients will show up on the right, and a list of recipes will be available for users to scroll through. To learn more about a specific recipe, users can press the ‘Explore’ button beneath the desired recipe.



Recipes displayed after inputting chicken, pepper, garlic, onion, and tomato.

Upon clicking into a recipe, detailed instructions will be shown, along with an ingredients list, serving size, prep time, wine-pairing, cuisine, and the source of the recipe. Let's click on the most liked recipe, ‘African Chicken Peanut Stew,’ and we can see a modal opens up as shown in the image below.

African Chicken Peanut Stew

Ready In : 45 minutes

Serving Size : 1

Cuisine : African

Wine Pairing : pinotage chenin blanc riesling

Link : Source

Ingredients :

Bell Peppers for garnishing

1.5 cups of Chopped Chicken

2.5 Cooking spoons of oil

1 teaspoon of Curry

2 garlic cloves

Small piece of Chopped ginger

1 cup of groundnut (Blended) or 1 Cooking spoon of peanut Butter

2 handfuls of Chopped onions

Pepper

Salt

Seasoning

1/2 small sweet potato (Chopped)

Pinch of thyme

1 Chopped small tomato

1.5 Cooking spoons of Blended tomato

Instructions :

- Season and Boil the Chicken for 10 minutes with salt, pepper, seasoning, a handful of onions. Once the chicken is ready, in the same stock, Boil the chopped sweet potatoes till its almost cooked. Save the stock in a separate Bowl and the chicken and potatoes in a separate Bowl as well. In a pot, heat up one cooking spoon of oil and fry the chicken till it Browns. Take it out and heat up the other 1.5 cooking spoons of oil and fry the onions, tomatoes Both chopped and Blended, ginger and garlic.
- Add your seasoning, curry, thyme, parsley, salt and pepper to the pot.
- Pour in your stock, chicken and potatoes to cook further. Stir in your peanut Butter and allow to cook for 10 minutes on low heat. If your sauce gets too thick, add a little water to it.
- Serve with white rice or more sweet potatoes. You could also garnish the dish with Bell peppers.

Recipe page with ingredients, instructions, prep time, serving size, cuisine-type, wine-pairing and a link to source.

Users accessing CookBook via the iOS or Android app will be able to directly upload images from their camera or their photo library, as shown in the images below.

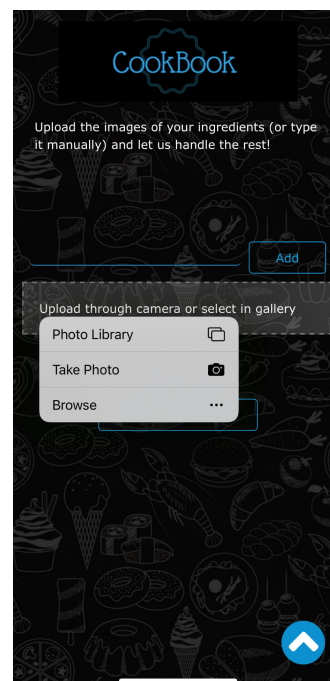



Image input on the Android app (left) and the iOS app (right).

Machine Learning Model

We used MobileNetV2 as the architecture for the machine learning model. MobileNets are characterized by their fast inference since they are based on a streamlined architecture, using depth-wise separable convolutions to build light weight deep neural networks [1]. This feature provides the basis for very efficient mobile-oriented models, allowing these models to serve as the base for many visual recognition tasks. Because we wanted a mobile version of our web app, we decided to use MobileNets. MobileNetV2 features an inverted residual structure with a bottleneck between layers, which also removes non-linearities in the narrow layers [2]. It also utilizes a novel framework known as SSDLite, which results in a prediction speed increase of up to 35% in comparison to previous versions [2].

We trained two separate implementations of a model on a dataset that consists of datasets from Kaggle (Fruits360, a dataset of images of fruits taken from all around it, and Fruit and Vegetable Image Recognition, a dataset include 36 species, each with around 100 images), and images from Bing Image Search and Google Image Search. After gathering all of the images, each image class was split into training and testing datasets in a ratio of 7:3 respectively. As training happened, we further increased the number of images by randomly applying data augmentations to them (random and rule cropping, flipping, rotation, adding noise using USM (unsharp masking), using histogram equalization, adjusting contrast and brightness). By doing so, we hoped to increase the accuracy. Finally, the model with the highest accuracy was picked for the final version of CookBook. The diagram below shows a list of the supported ingredients.

Apple	Clementine	Kaki	Nut Forest	Pomelo
Apricot	Cocos	Kiwi	Nut Pecan	Potato
Avocado	Corn	Kohlrabi	Onion	Quince
Banana	Corn Husk	Kumquats	Orange	Raddish
Beetroot	Cucumber	Lemon	Papaya	Rambutan
Blueberry	Dates	Lettuce	Passion Fruit	Raspberry
Cabbage	Eggplant	Limes	Pea	Redcurrant
Cactus fruit	Fig	Lychee	Peach	Salak
Cantaloupe	Garlic	Mandarine	Pear	Soy Beans
Capsicum	Ginger	Mango	Pepino	Spinach
Carambula	Granadilla	Mangostan	Pepper	Strawberry
Carrot	Grape	Maracuja	Physalis	Tamarillo
Cauliflower	Grapefruit	Melon Piel de Sapo	Physalis with Husk	Tangelo
Cherry	Guava	Mulberry	Pineapple	Tomato
Chestnut	Hazelnut	Napa Cabbage	Pitahaya Red	Turnip
Chilli Pepper	Huckleberry	Nectarine	Plum	Walnut
Choy	Jalepeno	Nectarine Flat	Pomegranate	Watermelon

Ingredients supported by the machine learning model.

Backend Flask

The backend has been implemented using Flask and is responsible for handling all the requests sent to it via the frontend. It listens on port 8080 and has 3 major functionalities — load the model.pht file and predict incoming images (:8080/recognise), create a spoonacular client instance to fetch lists of recipes using the ingredients (:8080/ingredients), and finally make a call to retrieve the complete info about a specific recipe (:8080/recipe). Since our ML model accepts one request at a time and predicts sequentially, multiple subsequent requests to the flask API slowed the system and in some cases even crashed or timed-out the server. Hence we came up with the idea to parallelize the task. We used docker and docker swarm to create multiple instances of the ML model with their own separate containers in the conda environment that could be scaled from 0 to even 1000 (needs more RAM on google compute engine as we keep on increasing dockers) at one click. This solved two major problems:

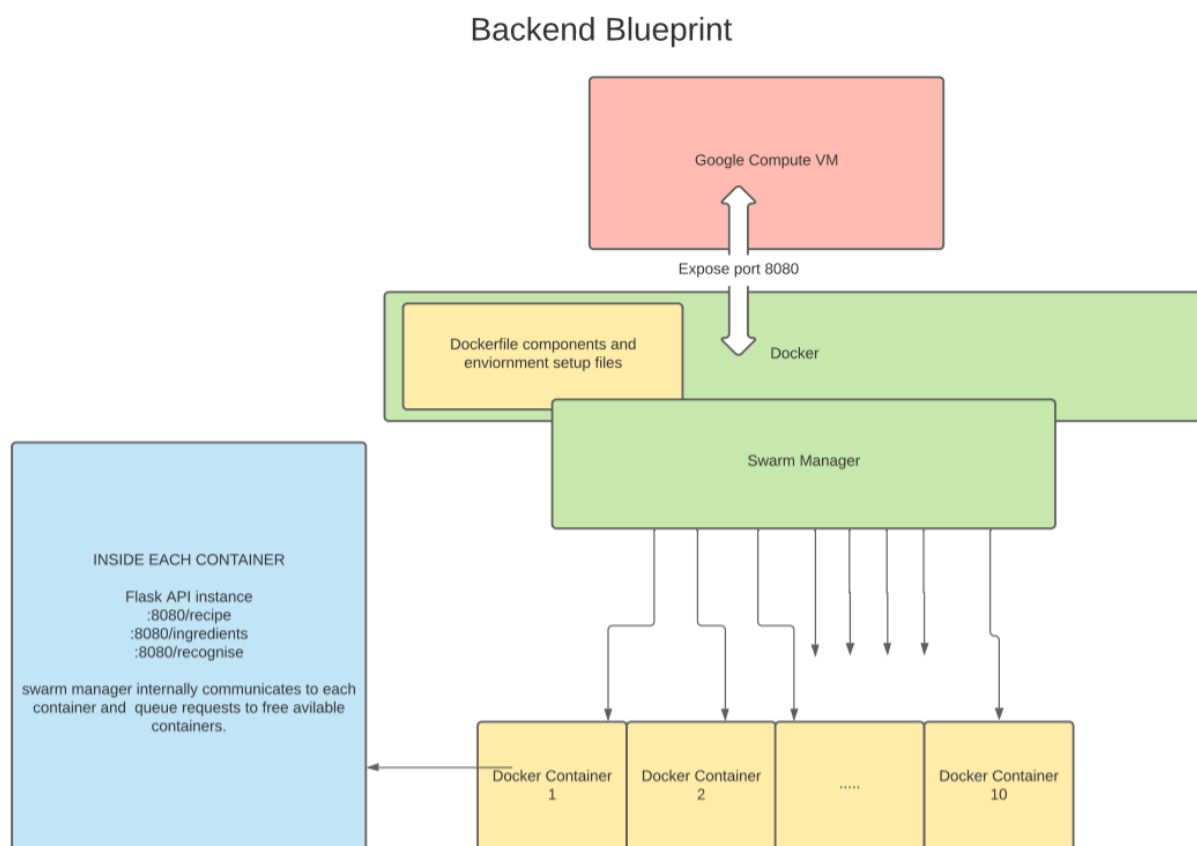
1. It can handle up to 10 simultaneous requests at once (current number of dockers as per optimal server configuration). This significantly reduces the time to respond to concurrent API calls by 90%.
2. Whenever a docker container crashes it doesn't affect Google's VM or other dockers. Instead, the manager node automatically transfers the workload to an available docker and in the meantime spins off a new docker with the same model instance to be used again. This reduces the downtime of our server by 99% in case of crashes or timeout error.

```
root@aryan-node: /home/baryan_edu/project-team-albertosaurus/reports/M4 - Google Chrome
ssh.cloud.google.com/projects/flash-rock-313815/zones/us-west1-b/instances/aryan-node?authuser=0&hl=en_US&projectNumber=604229437191&useAdminProxy=true
Every 2.0s: docker ps -a
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS              NAMES
6631de9da4c1      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.4.k3hvald478uddr2cqn8dt3eym
165c27b3174c      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.6.qh7qlarc0pbkoou7kledxmm6c
2e6add5a52d2      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.7.zyz7uwkka911ounkm7u0impyv
6542f2f6742a      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.10.fg4jvg4gg8xtorzgoh4bgo8hc
6b5312b6b99c      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.8.tt8bu7ax9oukxslufin619acw
60c55069b1c5      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.9.3mej29111430j2p8y0s3tqgaq
c0e16227fd29      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.1.jb0kqa4r4z9gw52q3bu34it5k
4f9f073deb75      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.2.jpwfdsy5o9tdxpujdzpm570i
33c6315eb7b3      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.3.taf7opg2jc6fnb30eri8mq6so
123047695ce3      cookbookimage:latest "conda run --no-capt..." 18 minutes ago      Up 18 minutes      8080/tcp           cookbook_flask.5.z80u09plw8pgdwoaiq78dzl0x
```

To build our app to be more robust and efficient, we thought of a way to increase the Spoonacular quota. As each authentication key is allowed for only 150 requests per day, each team member created a free account to total up our daily quota to 600 requests/day. Whenever

one of the key's quota is exhausted, the server automatically switches to the next available API key and begins retrieving data, thereby increasing our available resources by 4 times. The advantages of using a Spoonacular client in the Flask instead of directly making requests through the frontend was:

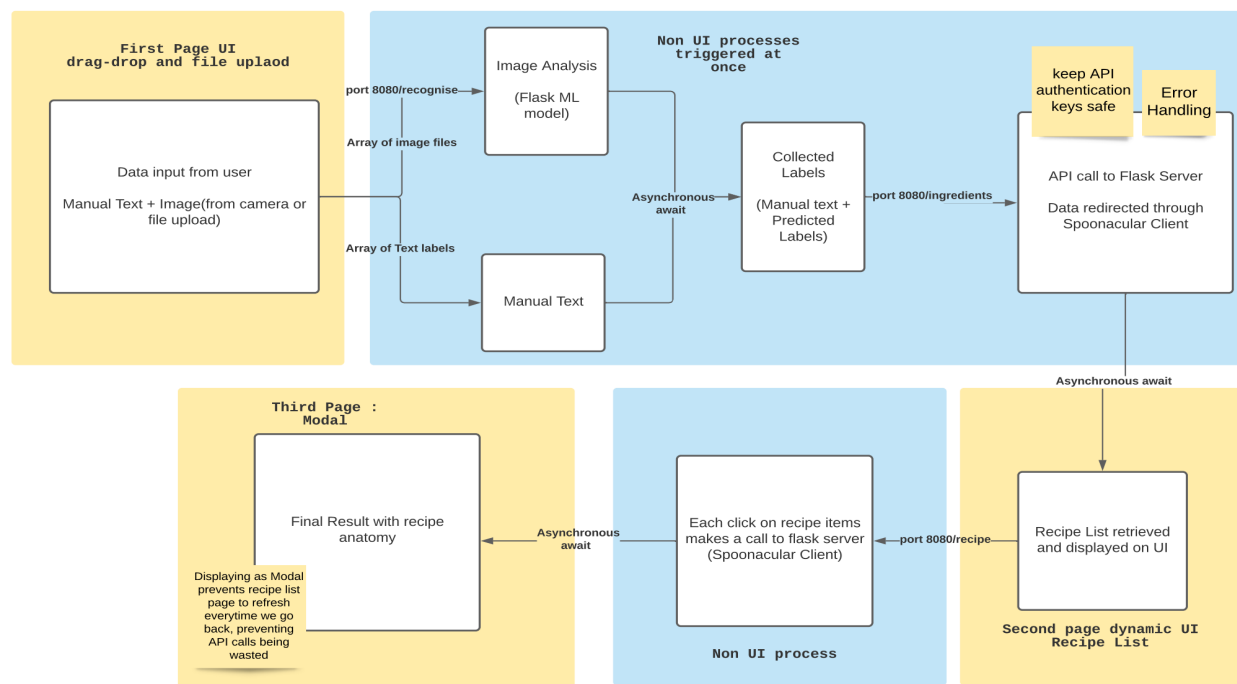
- Keep our authentication key safe and secure
- Easily handle switching
- Keep everything organised.



Web App

The web app was created using ReactJS and various npm libraries. Animations and page transitions have been used to improve user experience. All the components are functional components and use hooks to change and render states and the items are dynamically loaded using states and props. The complete structure of the frontend and how images get processed into

ingredients and further into recipes can be deduced from the blueprint [Fig 2]. Most tasks have been automated using multiple asynchronous calls and await functions. The image and manual feeds are segregated. The image files are uploaded to the server encoded in jsonified formdata through POST requests. The returned labels then combine with manual feed and a simultaneous request is made to flask for fetching recipes. The returned array of jsonified recipes are displayed dynamically as cards on the recipe list page each having a related explore button. Each 'explore' sends a POST request with an id to fetch info about the chosen recipe which is finally displayed as modal. The clever part about implementing modals is that they prevented refreshing of the recipe-list page whenever it was closed or open. As, fetching a list of 30 recipes at once was an expensive task regarding exhausting the spoonacular API limit, it helped reduce the cost by around 70%. The frontend has been deployed as a single instance on google cloud server and its default ip and port has been linked to a free domain name.



Android App

The Android app for CookBook can be installed by downloading CookBook.apk (found in M4 of the GitHub repository). In contrast to the web app, the Android app allows users to take photos directly from their camera and upload them. The app uses WebView to load the URL address of the web app and change the display to be more mobile-friendly. We first set the URL to the web

address with `webView.loadUrl (http://www.cookbookubc.ml)`, and then set up the `WebViewClient` and overrode the `shouldOverrideUrlLoading` method to ensure that the page would be loaded by the Android app instead of opening the web app on a browser. The next step was to set up the `WebChromeClient`, which intercepts HTML5 methods of calling to files and allows the Android operating system to handle the blocked files on its own. Then, the class `PopupWindow` was used to set up the popup layout and the popup events. To enable camera access, we had the app request for camera and photo album permissions. Additionally, to confirm whether an image was successfully obtained, a prompt message was implemented. If the fetching is successful, “Successful Access!” will be displayed. Otherwise, “Get failed” will be shown. Finally, the results were sent back to the web app in the form of an array of URIs.

iOS App

The iOS app for CookBook was successfully tested but cannot be installed by users due to restrictions on Apple developer accounts. The development process for the iOS app is similar to that of the Android app. We first set the fixed `NSURL` address as `http://www.cookbookubc.ml`, then built the `NSURLRequest` object and filled it in with the `NSURL` object. The third step was to initialize the `WKWebView` configuration and set up the layout. Then we called the `loadRequest` method of the `WKWebView` in the `viewDidLoad` method, which is the life cycle behavior declaration of `ViewController`, and filled it with the aforementioned `NSURLRequest` object.

Contributions

Jeremy created a machine learning model that was ultimately not picked for the final rendition of CookBook. He then focused on the secretarial side of things, such as the final report and the demo.

Aryan built the web app using ReactJS to analyze images with the embedded model in Flask and send and receive data using the Spoonacular API.

Ting created the final version of the machine learning model, embedded the trained model into Flask for the backend portion of the web app, completed both of the mobile apps, and helped explain her respective parts for the final report.

Rushil worked on an earlier implementation of the Android app and helped with the final report.

References

- [1] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam, "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications," 2017. [Online]. Available: [arXiv:1704.04861v1](https://arxiv.org/abs/1704.04861v1).
- [2] M. Sandler, A. Howard, M. Zhu, A. Zhmoginov, and L.-C. Chen, "MobileNetV2: Inverted Residuals and Linear Bottlenecks," 2019. [Online]. Available: [arXiv:1801.04381v4](https://arxiv.org/abs/1801.04381v4).